

# **EMPIRICAL ASSET PRICING VIA MACHINE LEARNING**

---

Lucas DUBOIS, Myriam LAMBORELLE, William LEMAIRE & Tom PARMENTIER

December 14, 2025

HEC Liège

# TABLE OF CONTENTS

INTRODUCTION:

PENALIZED REGRESSIONS:

DIMENSIONALITY REDUCTION:

NON-LINEAR MODELS & RESULTS

CONCLUSION:

## **INTRODUCTION:**

---

## OBJECTIVES:

The paper "Empirical Asset Pricing via Machine Learning" [2] has a twofold objective:

1. Setting new benchmarks to compare ML methods for asset pricing.
2. Synthesizing the asset pricing theory with the emerging ML literature.

## DEFINITION:

What is **Machine Learning (ML)**?

1. High-dimensional models.
  - ✓ Flexibility.
  - ✗ Overfitting.
2. “Regularization” methods for model selection.
  - ✓ Fix Overfitting.
3. Efficient algorithms for searching model specifications.
  - ✓ Fix non-linearities.

## DATASET AND METRICS:

- **Universe:** Monthly U.S. equities (CRSP/Compustat)
- **Sample:** 30,000 individual stocks.
- **Period:** 60 years (1957–2016)
- **Predictors:** 94 firm-level characteristics from the "factor zoo" + 74 industry-sector dummies + interaction terms.
  - 900+ predictors in total.

## BENCHMARK MODEL:

The authors use a characteristic OLS model (**OLS-3+H**) , similar to the famous Fama-French 3-factor model [1], as benchmark:

$$\hat{r}_{i,t+1} = \underbrace{\hat{\beta}_1 \text{SIZE}_{i,t} + \hat{\beta}_2 \text{VALUE}_{i,t} + \hat{\beta}_3 \text{MOM}_{i,t}}_{\text{3-factors}} + \hat{\delta}_t + \hat{\gamma}_{j(i)}. \quad (1)$$

This benchmark is commonly used in the empirical asset pricing literature, as well as in practice by asset managers.

## AN IMPORTANT METRIC:

Throughout the paper, the authors use the **out-of-sample  $R^2$**  to evaluate model performance:

$$R_{OOS}^2 = 1 - \frac{\sum_{t=T_0}^T (r_{t+1} - \hat{r}_{t+1})^2}{\sum_{t=T_0}^T (r_{t+1} - \bar{r})^2} \quad (1)$$

- Measures **raw predictive power** of a model.
- No **COMPLEXITY PENALTY** term.

## CROSS-VALIDATION FOR TUNING PARAMETERS:

The authors use a **rolling one-step-ahead out-of-sample cross-validation** procedure to select the **hyperparameters** of their models:

$$(\lambda^*, \rho^*) = \arg \min_{\lambda, \rho} \sum_{t=T_0}^{T-1} \sum_{i=1}^{N_t} (r_{i,t+1} - \hat{r}_{i,t+1}(\lambda, \rho))^2. \quad (2)$$

This allows to:

1. Account for **time-series dependence** in financial data.
2. Mimic how **investors** would forecast their portfolios.
3. Prevent **look-ahead bias** (terrible idea in finance!).

## PENALIZED REGRESSIONS:

---

## GENERAL IDEA:

*“The simple linear model is bound to fail in the presence of many predictors. When the number of predictors  $P$  approaches the number of observations  $T$ , the linear model becomes inefficient or even inconsistent.”*

Gu, Kelly & Xiu [2]

## WHY PENALIZE?

- Low **signal-to-noise ratio** in asset returns  $\Rightarrow$  overfitting risk!
- **Multicollinearity** among predictors  $\Rightarrow$  unstable estimates!

## INTRODUCING PENALTIES:

The basic framework for the penalized regressions takes place as a penalty in the **cost function**:

$$\mathbb{L}(\theta, \cdot) = \underbrace{\mathbb{L}(\theta)}_{\text{MSE}} + \phi(\theta, \cdot) \quad (3)$$

## ELASTIC-NET PENALTY:

The authors define the **Elastic-Net** penalty as:

$$\phi(\theta; \lambda, \rho) = \lambda(1 - \rho) \sum_{j=1}^P |\theta_j| + \frac{1}{2} \lambda \rho \sum_{j=1}^P \theta_j^2 \quad (4)$$

This notation can be converted to the one we used in class by letting  $\alpha = 1 - \rho$ .  
The tuning parameters are the **mixing parameter** ( $\rho$ ) and the **strength parameter** ( $\lambda$ ).

We can find our *special cases*:

$$\begin{cases} \text{Lasso: } \rho = 0 \\ \text{Ridge: } \rho = 1 \end{cases}$$

## OPTIMIZATION ALGORITHM:

---

### Algorithm 1 Accelerated Proximal Gradient Method (FISTA)

---

**Initialization:**  $\theta_0 = 0$ ,  $m = 0$ , step size  $\gamma$

**while**  $\theta_m$  not converged **do**

$$\bar{\theta} \leftarrow \theta_m - \gamma \nabla L(\theta) \Big|_{\theta=\theta_m}$$

$$\tilde{\theta} \leftarrow \text{prox}_{\gamma\phi}(\bar{\theta})$$

$$\theta_{m+1} \leftarrow \tilde{\theta} + \frac{m}{m+3}(\tilde{\theta} - \theta_m)$$

$$m \leftarrow m + 1$$

**end while**

**Result:** Final parameter estimate  $\theta_m$

---

## DIMENSIONALITY REDUCTION:

---

## BASIC FRAMEWORK:

$$R = Z\theta + E, \quad (5)$$

where  $\begin{cases} R \in \mathbb{R}^{NT \times 1} & \text{is the stacked vector of returns } r_{i,t+1}, \\ Z \in \mathbb{R}^{NT \times P} & \text{is the stacked predictor matrix } z'_{i,t}, \\ \theta \in \mathbb{R}^{P \times 1} & \text{is the full coefficient vector,} \\ E \in \mathbb{R}^{NT \times 1} & \text{is the stacked residual vector } \varepsilon_{i,t+1}. \end{cases}$

## BASIC FRAMEWORK:

$$R = (\mathbb{Z}\Omega_K) \theta_K + \tilde{E} \quad (6)$$

where  $\begin{cases} \Omega_K \in \mathbb{R}^{P \times K} & \text{contains the component weights } (w_1, \dots, w_K), \\ \mathbb{Z}\Omega_K \in \mathbb{R}^{NT \times K} & \text{is the dimension-reduced predictor matrix,} \\ \theta_K \in \mathbb{R}^{K \times 1} & \text{is the coefficient vector in reduced space,} \\ \tilde{E} \in \mathbb{R}^{NT \times 1} & \text{is the new residual vector.} \end{cases}$

## GENERAL IDEA:

- $P$  can be very large (hundreds or thousands of characteristics), but expected returns live in a *low-dimensional* space.
- PCR / PLS construct  $K \ll P$  predictive components:

$$\textcolor{blue}{Z}\Omega_K \in \mathbb{R}^{NT \times K}, \quad \Omega_K = [w_1, \dots, w_k].$$

- Each  $w_j$  is a set of weights forming a linear combination of characteristics (a “latent factor” or component).
- Instead of forecasting with many noisy, collinear signals, we forecast with a few stable, informative components that summarize the underlying economic forces driving returns.

## PENALIZED REGRESSION VS. DIMENSION REDUCTION

	Penalized Regression	Dimension Reduction
Goal	shrink coefficients $\theta$	reduce predictors to $K \ll P$
Form	$\min_{\theta} L(\theta) + \phi(\theta)$	$R = (Z\Omega_K)\theta_K + \tilde{E}$
Dimensionality	stays at $P$	becomes $K$
Complexity Control	penalty (L1/L2/EN)	projection $Z \rightarrow Z\Omega_K$
Examples	Ridge, Lasso, EN	PCR, PLS

## PRINCIPAL COMPONENTS REGRESSION (PCR)

1. Compute eigenvectors of  $Z^T Z$ .
2. Select the top  $K$  eigenvectors:

$$\Omega_K = [w_1, \dots, w_K].$$

3. Construct components:

$$C = Z\Omega_K.$$

4. Run OLS on reduced predictors:

$$R = C\theta_K + \tilde{E}.$$

PCR finds directions in  $Z$  with the *largest variance*. It assumes that predictive power lies in high-variance directions of the characteristic space.

## PARTIAL LEAST SQUARES (PLS)

1. First component solves:

$$w_1 = \arg \max_{\|w\|=1} (\mathbf{R}^T \mathbf{Z} w)^2.$$

2. Form the first score:

$$c_1 = \mathbf{Z} w_1.$$

3. Deflate  $\mathbf{Z}$  and  $\mathbf{R}$ ; repeat to obtain  $w_2, \dots, w_K$ :

$$\Omega_K = [w_1, \dots, w_K].$$

4. Regress on reduced predictors:

$$\mathbf{R} = (\mathbf{Z} \Omega_K) \theta_K + \tilde{\mathbf{E}}.$$

PLS chooses directions *most useful for predicting returns*, making it supervised and more targeted than PCR.

## PCR vs. PLS

	PCR	PLS
Optimizes	variance of $Z$	covariance with $R$
Supervised?	No	Yes
Components	eigenvectors of $Z^T Z$	directions maximizing $R^T Z w$
Reduced predictors	Variance components	Predictive components
Strength	Captures structure of $Z$	Return-predictive signals
Weakness	Ignores $R$ in components	Nonlinearities(?)

## **NON-LINEAR MODELS & RESULTS**

---

## WHY GO NON-LINEAR?

### The Linear Limitation:

- Linear models assume the effect of a predictor (e.g., Value) is constant, regardless of other conditions (e.g., Volatility).
- **Reality:** Asset pricing is likely conditional.
  - *Example:* "Value" might only work well during high volatility.

### The Machine Learning Solution:

- Use flexible functional forms  $g(z_{i,t})$  to capture **interactions** without specifying them ex-ante.
- We examine three classes: **GLM**, **Regression Trees**, and **Neural Networks**.

## 1. GENERALIZED LINEAR MODELS (GLM)

GLM introduces non-linearity via **Splines** (piecewise polynomials) for each predictor individually.

$$g(z; \theta) = \sum_{j=1}^P \sum_{k=1}^K \theta_{j,k} p_k(z_j) \quad (7)$$

- **Penalty:** Uses **Group Lasso** to select or drop entire characteristics (group of spline terms) at once.
- **Limitation:** It is additive. It captures non-linearities of *single* variables ( $z_j^2$ ), but **misses cross-variable interactions** ( $z_i \times z_j$ ).

## 2. TREE-BASED METHODS

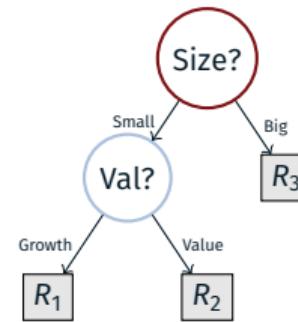
Trees partition the predictor space into rectangular regions and predict the average return in each "leaf."

### Random Forests (RF):

- Builds  $B$  deep, decorrelated trees.
- Uses **Dropout**: Randomly subsets predictors at each split to force diversity.
- Prediction is the *average* of all trees (reduces variance).

### Gradient Boosted Trees (GBRT):

- Sequentially builds shallow "weak learners."
- Each tree corrects the errors of the previous one.



*Captures interactions naturally.*

### 3. NEURAL NETWORKS (NN)

The most flexible model. Composed of an **Input Layer**, **Hidden Layers** (nonlinearities), and an **Output Layer**.

#### Feed-Forward Architecture

For a neuron  $k$  in layer  $l$ , the value is a non-linear transformation of the previous layer:

- **Activation:** Rectified Linear Unit ( $\text{ReLU}(x) = \max(0, x)$ ).
- **Regularization:** Essential to prevent overfitting.
  - $L_1$  Penalty, Early Stopping, Batch Normalization, Ensembles.
- **Architectures:** Tested NN1 (1 layer) to NN5 (5 layers).

## RESULTS: OUT-OF-SAMPLE $R^2$

*Non-linear models dominate linear benchmarks.*

Model	OLS-3	En-Net	GLM	RF	NN3
$R^2_{OOS}$ (%)	0.16	0.11	0.19	0.33	<b>0.40</b>

- **GLM Failure:** Performance is close to linear models  $\Rightarrow$  Univariate non-linearity is *not* enough.
- **NN3 Peak:** Performance peaks at 3 layers. "Shallow" learning works best for finance (low signal-to-noise).
- **Economic Value:** Small  $R^2$  gains translate to large profits.

## RESULTS: ECONOMIC GAINS (SHARPE RATIOS)

Annualized Sharpe Ratios for Long-Short Decile Portfolios (Value-Weighted):

- Neural Networks **more than double** the risk-adjusted returns of the OLS benchmark.

## WHAT MATTERS? (INTERPRETABILITY)

By analyzing variable importance (Sensitivity Analysis), models agree on the dominant signals:

1. **Price Trends:** Momentum (12m), Short-term Reversal (1m).
2. **Liquidity:** Market Cap, Dollar Volume, Bid-Ask Spread.
3. **Volatility:** Idiosyncratic Volatility, Beta.

*Why do NN and Trees win? They capture **Interaction Effects**.*

- *Example:* The "Reversal" effect is strong for small stocks but weak/concave for large stocks. Linear models miss this nuance.

## **CONCLUSION:**

---

## CONCLUSION:

What the Paper brings to **Empirical Asset Pricing**

1. Systematic Comparison
2. Bridges tradition with modernity
3. Shows limits

## CONCLUSION:

### The Dataset

1. 30 000+ U.S. stocks
2. Almost 60 years of monthly data
3. 900+ predictors from the "factor zoo"
4. Traditional models

$R^2_{oos}$  **BEST MODELS:**

	NN3	NN4	NN2	NN5	RF	GBRT+H
All	0.40	0.39	0.39	0.36	0.33	0.34
Top 1,000	0.67	0.62	0.70	0.64	0.63	0.52
Bottom 1,000	0.47	0.46	0.45	0.42	0.35	0.32

$R^2_{oos}$  WORST MODELS:

	OLS+H	ENet+H	OLS-3+H	GLM+H	PCR	PLS
All	-3.46	0.11	0.16	0.19	0.26	0.27
Top 1,000	-11.28	0.25	0.31	0.14	0.06	-0.14
Bottom 1,000	-1.30	0.20	0.17	0.30	0.34	0.42

## SHARPE RATIO (H-L) BEST MODELS:

	NN4	NN3	NN1	NN2	NN5	RF
SR(H-L)	1.35	1.20	1.17	1.16	1.15	0.98

## SHARPE RATIO (H-L) WORST MODELS:

	ENet+H	OLS-3+H	PLS	GLM+H	GBRT+H	PCR
SR(H-L)	0.39	0.61	0.72	0.76	0.81	0.88

## CONCLUSION:

### Key Takeaways

1. Return prediction improvement
2. Asset pricing evolution

### Limitations of the Paper

1. Interpretability
2. Large datasets
3. International robustness

### Final Message

## FINAL MESSAGE:

### The Dataset

1. New approach
2. Important modernization step

## REFERENCES:

- [1] Eugene F. Fama and Kenneth R. French. “**Common risk factors in the returns on stocks and bonds**”. en. In: *Journal of Financial Economics* 33.1 (Feb. 1993), pp. 3–56. ISSN: 0304405X. DOI: [10.1016/0304-405X\(93\)90023-5](https://doi.org/10.1016/0304-405X(93)90023-5). URL: <https://linkinghub.elsevier.com/retrieve/pii/0304405X93900235> (visited on 11/21/2025).
- [2] Shihao Gu, Bryan Kelly, and Dacheng Xiu. “**Empirical Asset Pricing via Machine Learning**”. en. In: *The Review of Financial Studies* 33.5 (May 2020). Ed. by Andrew Karolyi, pp. 2223–2273. ISSN: 0893-9454, 1465-7368. DOI: [10.1093/rfs/hhaa009](https://doi.org/10.1093/rfs/hhaa009). URL: <https://academic.oup.com/rfs/article/33/5/2223/5758276> (visited on 10/24/2025).

